Data Profiling SideCar

Data Needs Observability

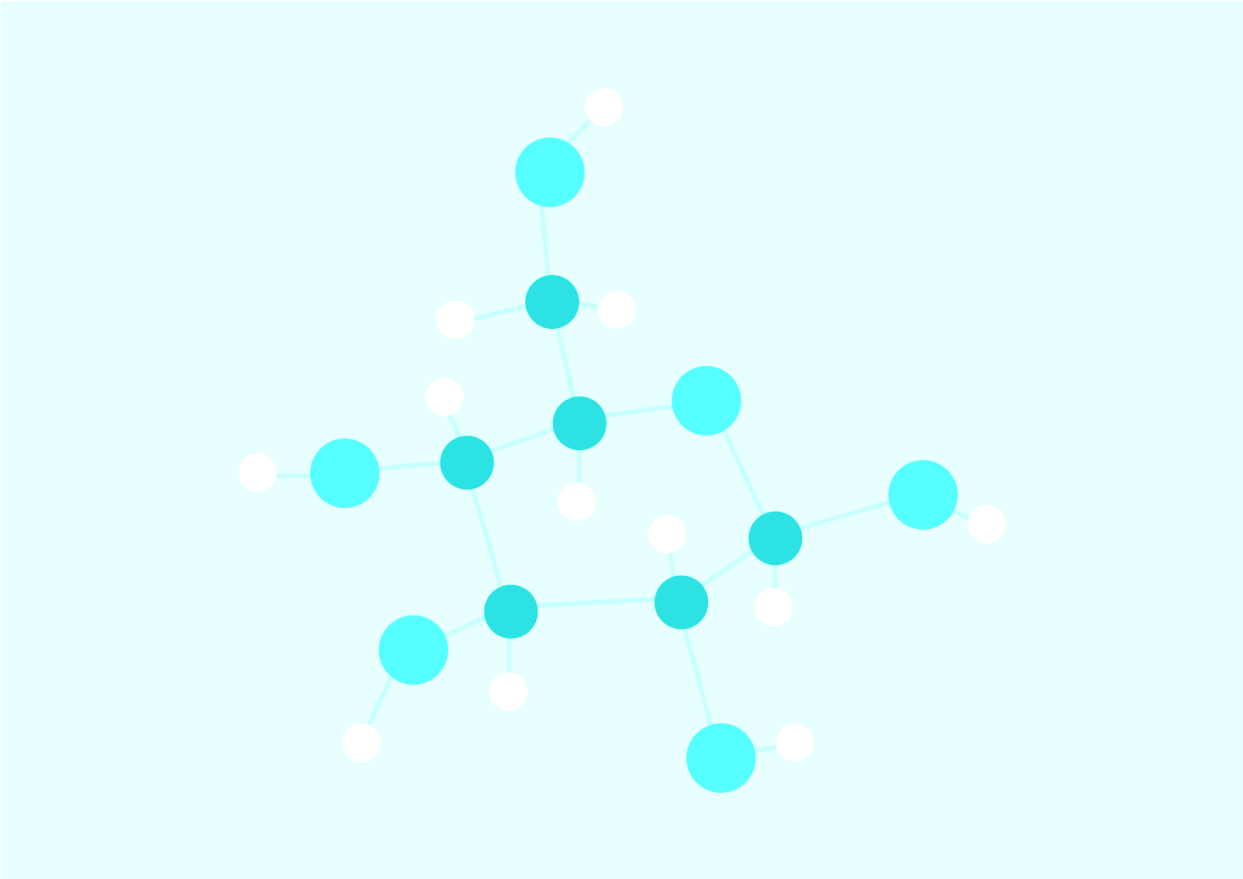

Aniruddha Mitra

Field CDAO

Google CLoud Platform

mitracdao@gmail.com | 240-728-4667

Executive Summary

In modern financial data ecosystems, the greatest threat is not the "hard failure" of a pipeline, but Silent Corruption—data that passes every technical schema check yet is statistically mutated at the population level. Traditional inline data quality tools often fail to detect these shifts without significant latency penalties. This paper proposes the Value Profiling Sidecar, a decoupled, AI-native architecture that leverages statistical process control (SPC) and semantic fingerprinting to provide zero-latency, lineage-aware observability.

Table of Contents

Executive Summary 1
Part 1 — What value profiling actually measures 2
Family 1: Univariate distribution statistics 2
Family 2: Distribution distance metrics 3
Family 3: Statistical Process Control 4
Family 4: Anomaly detection 5
Part 2 — Cross-column and temporal profiling 6
Cross-column profiling 6
Ratio invariant monitoring 6
Referential completeness 6
Temporal profiling 6
Seasonality-adjusted anomaly detection 7
Solutions ARchitecture: 7
Architecture 1 — Inline Pipeline Profiling 8
Architecture 2 — Dedicated Profiling Sidecar 9
Architecture 3 — Full Observability Fabric 10
The semantic fingerprint in detail 11
The Profiling Sidecar: Architecture and Design in Detail 12
The core design principle: observe, don't intercept 12
What a profile event contains 13
The five internal engines 14
Engine 1 — Univariate engine 14
Engine 2 — Cross-column engine 15
Engine 3 — Temporal engine. 15
Engine 4 — ML anomaly engine 15
Engine 5 — Fingerprint engine 16
The baseline store: the sidecar's memory 16
Signal correlation: turning multiple alerts into one incident 18
The lineage fan-out: why proactive notification matters 19
Operational characteristics and deployment considerations 20
Conclusion: **21**

Introduction:

In the high-stakes environment of modern financial services, the most dangerous threat to an enterprise is not the pipeline that breaks, but the pipeline that appears to work perfectly while delivering "silent corruption."

As organizations migrate from legacy on-premises environments to AI-native cloud architectures, they encounter a critical paradox: while data volume and velocity have scaled exponentially, our ability to verify the statistical integrity of that data remains tethered to manual, high-latency processes. Traditional data quality tools are designed for "hard failures"—detecting null values or schema mismatches. They are largely blind to "soft mutations," where data remains technically valid but becomes statistically fraudulent, leading to catastrophic errors in regulatory reporting, risk modeling, and automated financial decision-making.

The challenge is further compounded by the "latency tax." In a traditional architecture, performing deep statistical profiling within the critical data path creates a bottleneck, forcing leaders to make an impossible choice: deliver insights at the speed the business demands, or slow down to ensure those insights are trustworthy.

This white paper introduces a new architectural standard for the "Cognitive Enterprise": the **Value Profiling Sidecar**. By decoupling complex statistical validation from the execution layer, we enable a zero-latency observability fabric that proactively identifies mutations before they reach the consumer. This is the foundational technology required to move beyond reactive data management and into a state of **Data Nirvana**—where data is not just a managed asset, but a "bankable" engine of growth and autonomous intelligence.

What value profiling actually measures

Value profiling is the practice of building a statistical model of what a column's values look like under normal conditions, then testing every new batch of data against that model. The key insight is that a single value check ("is this number in range?") is almost always insufficient. What matters is whether the *distribution* of values has shifted —

which can happen even when every individual value is technically valid.

There are four families of statistical methods, each catching a different class of mutation.

Family 1: Univariate distribution statistics

These describe the shape of a single column's value distribution. They are the first line of detection.

For numeric columns the core statistics are mean, median, standard deviation, skewness, kurtosis, and the full percentile profile (p1, p5, p10, p25, p50, p75, p90, p95, p99, p99.9). The percentile profile matters more than mean/stddev for financial data because financial distributions are almost never normal — they have fat tails, and changes at the tail are often where mutations hide. A rogue transaction that doubles the p99 value while leaving the mean unchanged is invisible to a mean-based check and immediately visible in a percentile profile comparison.

For categorical columns the analogous statistics are value frequency distribution (the empirical probability mass function), number of distinct values, and entropy. A lookup table reclassification that shifts 8% of records from category A to category B shows up immediately as a change in the frequency distribution, even if every individual value is a valid category.

For date/timestamp columns: min, max, gaps in sequence, distribution of records by day-of-week and hour-of-day (which should be stable for most operational data), and the rate of change over time.

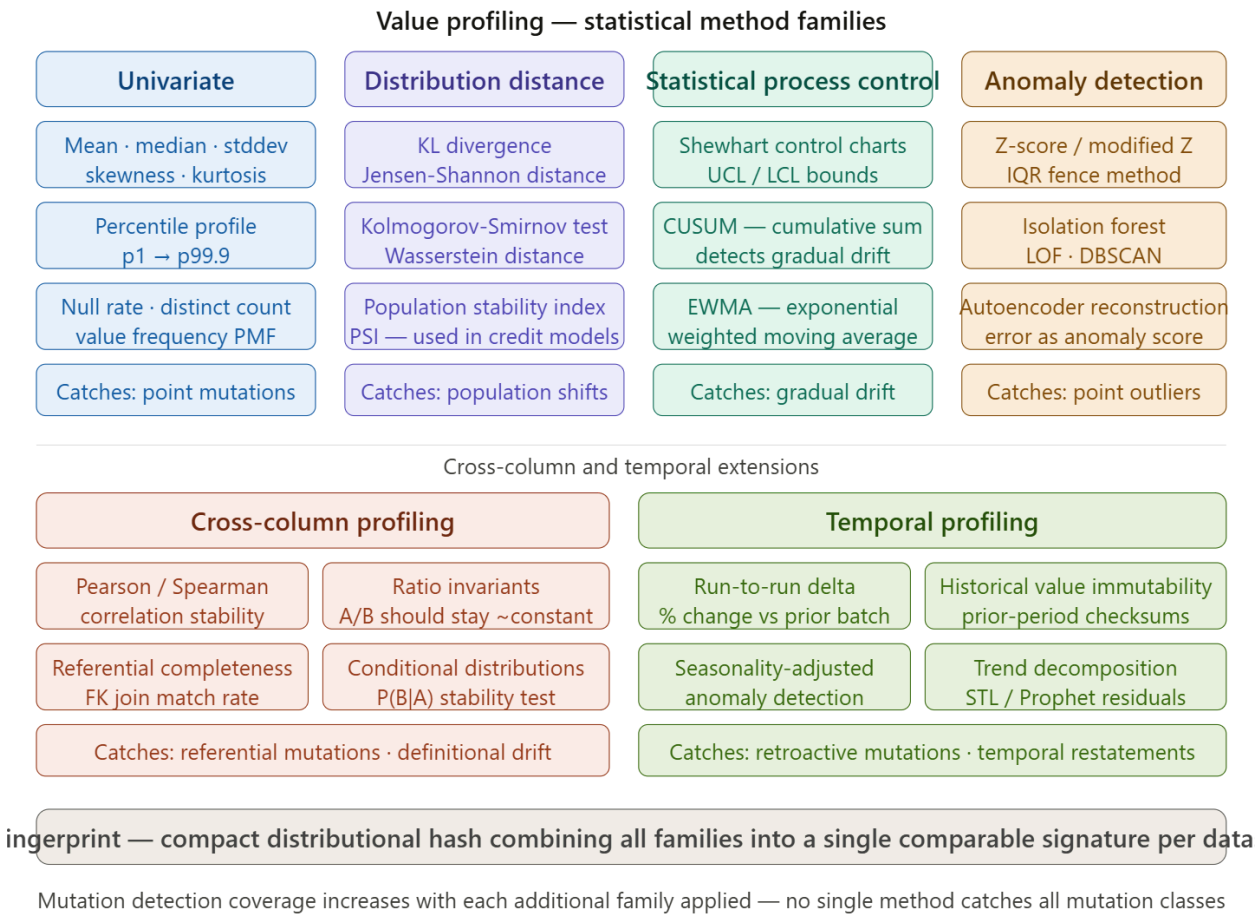

Family 2: Distribution distance metrics

Where univariate statistics describe a single snapshot, distribution distance metrics compare two snapshots — this run versus the reference baseline — and quantify how different they are. These are the workhorses of mutation detection at population scale.

The **Kolmogorov-Smirnov (KS) test** is the most commonly used. It computes the maximum absolute difference between two empirical cumulative distribution functions. A KS statistic above a calibrated threshold triggers an alert. It is non-parametric (makes no assumptions about the distribution shape), has well-understood statistical properties, and is computationally cheap. Its weakness: it is most sensitive to differences in the middle of the distribution and less sensitive to tail changes.

The **Population Stability Index (PSI)** is widely used in credit risk modeling and increasingly applied to pipeline monitoring. PSI buckets the distribution into deciles,

computes the percentage of records in each bucket for the baseline and current run, and calculates a weighted divergence. A PSI above 0.2 conventionally triggers a "significant shift" alert. PSI has the advantage of being interpretable — you can see which buckets drove the shift — and is well-understood by risk teams in financial services.

The **Wasserstein distance** (also called Earth Mover's Distance) measures how much "work" it would take to transform one distribution into the other. It is more sensitive to changes in distribution shape than KS and more robust to outliers than simple difference metrics. It is computationally more expensive but scales well on modern data platforms with approximate algorithms.

KL divergence and Jensen-Shannon distance are information-theoretic measures. KL divergence is asymmetric ($D(P||Q) \neq D(Q||P)$), which makes it sensitive to the direction of the shift. JS distance is the symmetric version. Both are useful for categorical columns where you want to measure how much the probability mass has moved between value categories.

Family 3: Statistical Process Control

SPC methods were developed in manufacturing to distinguish between natural process variation (which should be tolerated) and special cause variation (which indicates something has changed and needs investigation). They are particularly useful for financial data pipelines because financial metrics have natural seasonality and volatility that makes simple threshold-based alerts noisy.

A **Shewhart control chart** establishes upper and lower control limits (UCL/LCL) at $\pm 3\sigma$ from the process mean, computed from a historical baseline. Any observation outside those limits triggers an alert. The rule of thumb from manufacturing: a single point outside the 3σ bounds is a strong signal; eight consecutive points on the same side of the mean (without crossing it) is also a strong signal of sustained drift. Both rules are directly applicable to pipeline metric monitoring.

CUSUM (Cumulative Sum) is specifically designed to detect small, sustained shifts in a process mean — the class of change that Shewhart charts miss. A CUSUM chart accumulates the deviation from the target mean over time; when the cumulative sum crosses a threshold, an alert fires. For financial pipelines where mutations often produce small systematic shifts (a 2% change in all values in a category, a 50 basis point shift in a

rate column), CUSUM is dramatically more sensitive than a Shewhart chart.

EWMA (Exponentially Weighted Moving Average) applies exponentially decaying weights to historical observations — recent observations matter more than older ones. This makes it responsive to gradual drift while being robust to point outliers. EWMA is the right method when the concern is slow definitional drift rather than sudden point mutations.

Family 4: Anomaly detection

For detecting individual records that are anomalous relative to the population — the most visible class of mutation.

The **Z-score** is the simplest: `z = (x - μ) / σ`. Values with |z| > 3 are flagged. The modified Z-score replaces mean and stddev with median and MAD (Median Absolute Deviation), making it robust to the contamination effect — where a large mutation inflates the mean and stddev, making itself harder to detect.

The **IQR fence method** flags values below `Q1 - 1.5 × IQR` or above `Q3 + 1.5 × IQR` (Tukey's method). It makes no assumptions about normality and is the right choice for heavily skewed financial distributions.

Isolation Forest is the leading ML-based anomaly detection method for tabular data. It works by repeatedly partitioning the feature space with random splits; anomalous points are isolated quickly (require fewer splits) compared to normal points. It handles high-dimensional feature spaces (multiple columns simultaneously), scales to large datasets, and provides an anomaly score rather than a binary flag — which enables materiality-weighted alerting.

Autoencoders are the most powerful option for complex, high-dimensional data. Train an autoencoder on clean historical data; the reconstruction error for a new batch is the anomaly score. Records with high reconstruction error are anomalous. The advantage over Isolation Forest: the autoencoder learns the full distributional structure of the data, including complex non-linear relationships between columns. The cost: requires a training dataset, takes more infrastructure to operate, and is harder to explain to a regulator.

Cross-column and temporal profiling

Cross-column profiling

It is where most production mutation detection gaps live. A pipeline can pass all univariate checks on every column individually and still be silently corrupted if the *relationship between columns* has changed.

The key methods: **correlation stability testing** computes the Pearson or Spearman correlation between column pairs on the current batch and compares it against the historical baseline correlation. A sudden drop in correlation between `interest_income` and `average_earning_assets` (which should be highly correlated) is a stronger signal than any check on either column individually.

Ratio invariant monitoring

It is particularly powerful for financial data. Many financial relationships should be approximately constant: assets = liabilities + equity (the balance sheet identity), NIM = interest income / average earning assets (a business invariant), total rows = sum of segment rows (an additive decomposition invariant). These ratios should be stable within a narrow band across every pipeline run. A ratio that moves outside that band signals either a mutation or a definitional change — both of which warrant investigation.

Referential completeness

It measures the join match rate between fact tables and dimension tables. If 98% of `account_id` values in the fact table join to the dimension table normally and that rate drops to 91%, a dimension table change has orphaned 7% of facts. This is precisely the referential mutation pattern that causes reports to silently drop records.

Temporal profiling

It addresses the retroactive restatement problem — the hardest class of mutation to detect because the pipeline runs normally; it is only historical data that has changed. The solution is **historical value immutability checksums**: at the end of each pipeline run, compute a deterministic checksum of all prior-period values that should be immutable. On subsequent runs, recompute those checksums and compare. Any discrepancy indicates a retroactive change. This requires storing checksums (cheap) not re-reading all

historical data (expensive).

Seasonality-adjusted anomaly detection

It is essential for financial data. A 15% spike in transaction volume on the last day of the quarter is not a mutation — it is quarter-end behavior. A 15% spike on a random Wednesday in February is. Naive anomaly detection without seasonality adjustment produces enormous false positive rates on financial data. The right approach uses STL decomposition (Seasonal-Trend decomposition using Loess) or Facebook Prophet to separate seasonal, trend, and residual components, then applies anomaly detection to the residual.

Solutions Architecture:

Here are the three solution architectures — showing how these methods are assembled into deployable systems at different scales of sophistication.

Architecture 1 is the in-pipeline inline layer — the minimum viable profiling system that can be built into any existing orchestration framework today.

Architecture 2 is the dedicated profiling sidecar — a purpose-built statistical monitoring service that runs alongside the pipeline without being embedded in it.

Architecture 3 is the full observability fabric with lineage-aware propagation — the complete end-state described in the research paper.

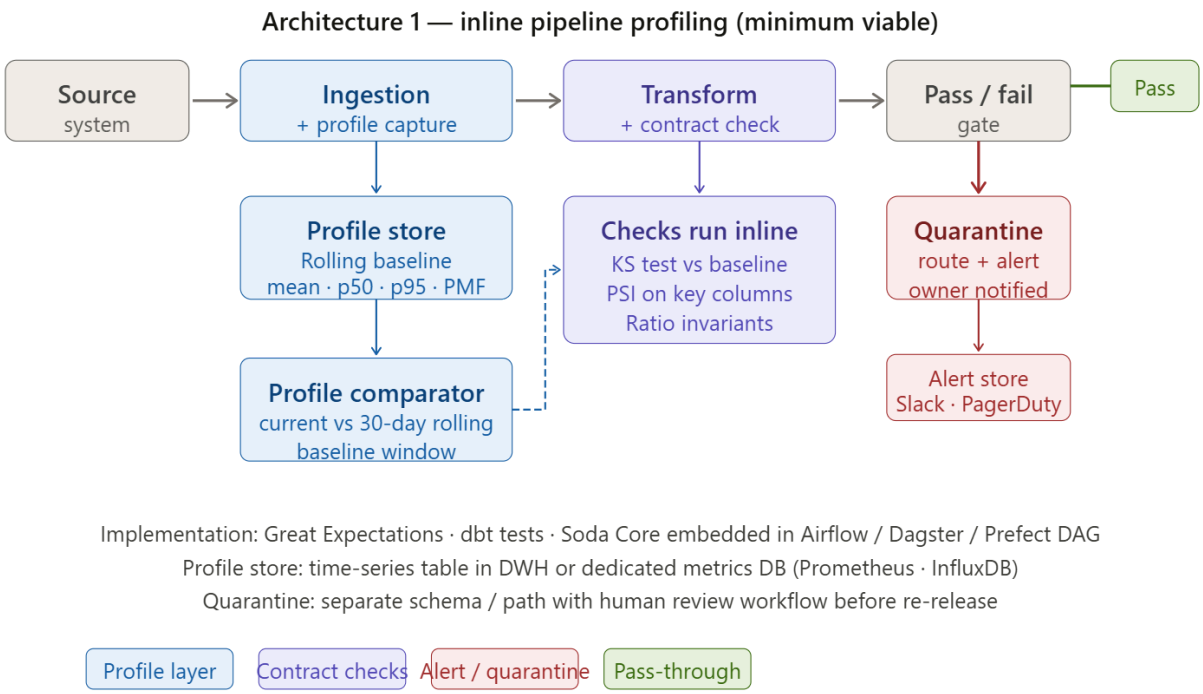

Architecture 1 — Inline Pipeline Profiling

It is the starting point for any team. The profiling logic is embedded directly in the pipeline DAG as pre- and post-transformation steps. At ingestion, statistical profiles are captured and stored in a rolling baseline table. At transformation, contract checks run inline comparing the current batch against the baseline. If checks pass, data proceeds. If they fail, data is routed to quarantine and an alert fires. Implementation tools: Great Expectations or Soda Core embedded in Airflow, Dagster, or Prefect. Profile storage: a time-series table in the data warehouse or a lightweight metrics database. This architecture adds 5–15% pipeline overhead and catches the majority of point mutations and population shifts at source. Its limitation: it has no view of the downstream consequence of a detected mutation.

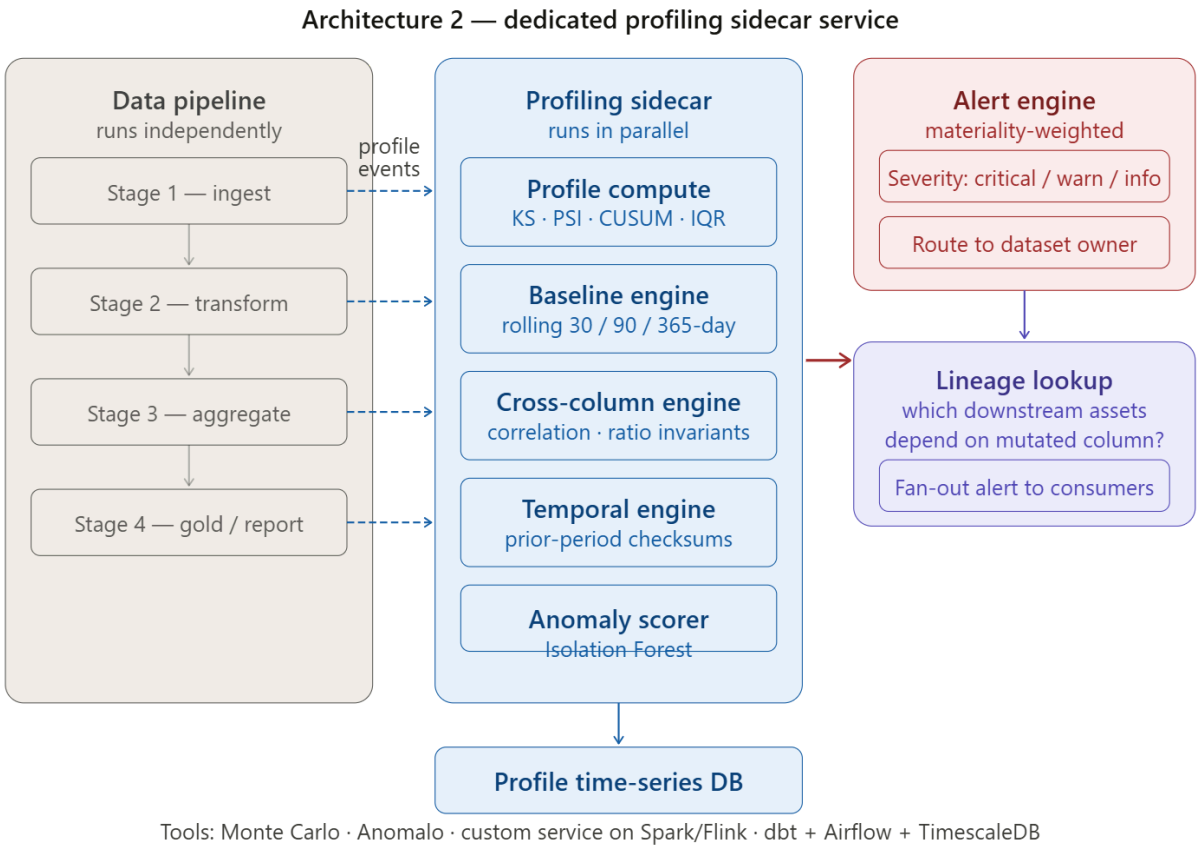

Architecture 2 — Dedicated Profiling Sidecar

It decouples the profiling logic from the pipeline execution. The pipeline runs normally and emits profile events (statistical snapshots of each dataset at each stage boundary) onto a metadata bus. A separate profiling service consumes those events, runs the full battery of statistical checks — KS tests, PSI, CUSUM, cross-column correlation, prior-period checksums, Isolation Forest scoring — and routes anomalies to an alert engine that applies materiality weighting and lineage-based fan-out. The key advantage over Architecture 1: the profiling service can run more computationally intensive checks (including ML-based anomaly detection) without adding to the pipeline's critical path latency. It also maintains a rich historical profile time series that enables trend analysis and threshold auto-calibration. Implementation options: dedicated commercial tools (Monte Carlo, Anomalo, Bigeye) or a custom service built on Spark Structured Streaming consuming OpenLineage events, with anomaly scoring in Python and alerts via PagerDuty.

Now the full end-state architecture — the complete observability fabric with

lineage-aware propagation and automated impact assessment.

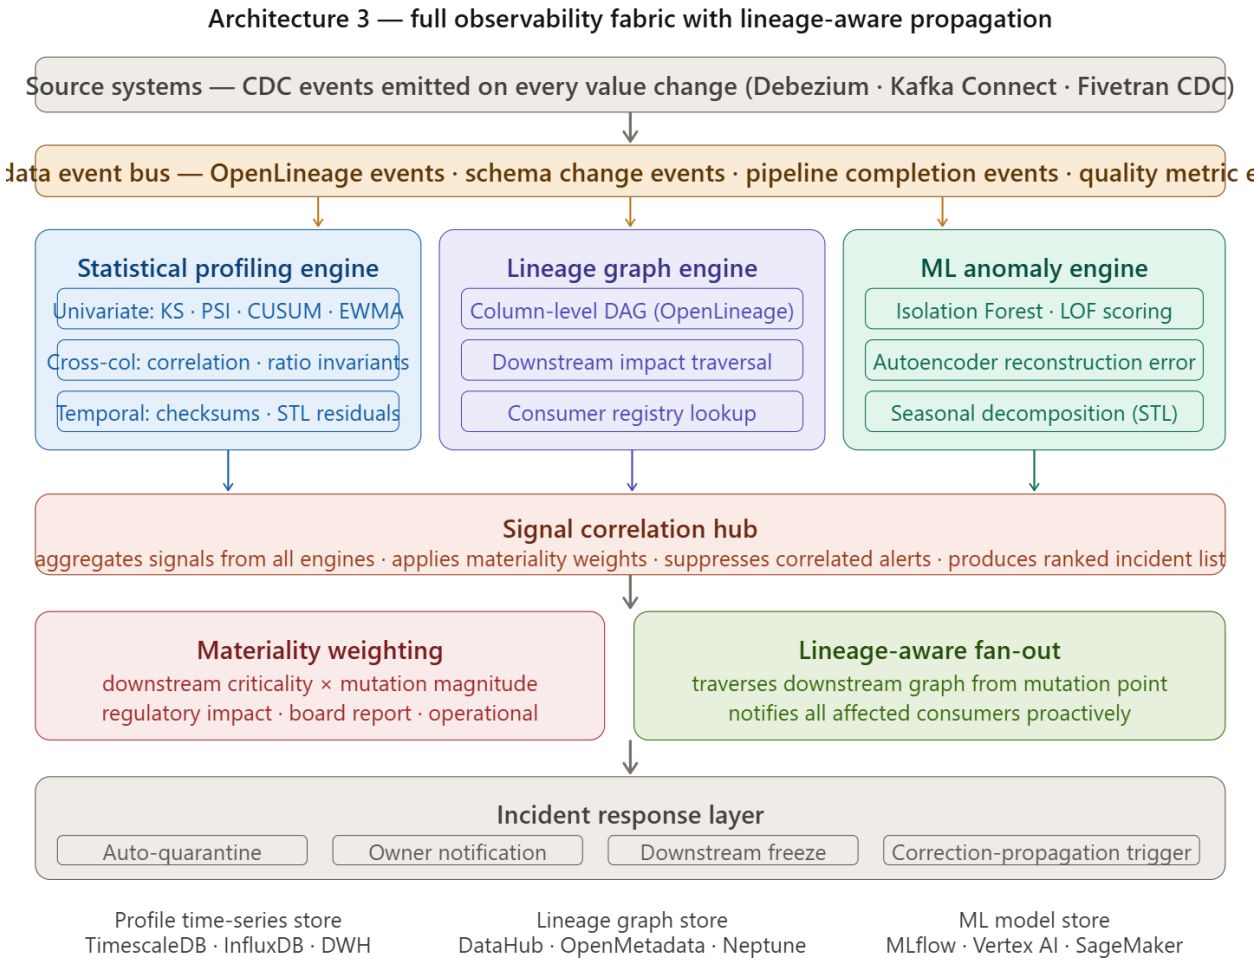

Architecture 3 — Full Observability Fabric

It is the complete end-state. Three parallel engines consume the metadata event bus simultaneously: the statistical profiling engine (running the full battery of univariate, cross-column, and temporal checks), the lineage graph engine (maintaining the column-level dependency graph and traversing it on mutation detection), and the ML anomaly engine (running Isolation Forest and autoencoder scoring for high-value datasets). Their outputs are aggregated in a signal correlation hub that suppresses correlated alerts, applies materiality weights, and produces a ranked incident list. The materiality engine assigns severity based on downstream criticality — a mutation feeding a regulatory capital ratio gets critical severity; one feeding an internal research report gets informational. The lineage-aware fan-out traverses the downstream graph

from the mutation point and notifies all affected consumers proactively, before their next scheduled pipeline run. The incident response layer handles auto-quarantine, owner notification, downstream pipeline freeze, and correction-propagation — the last ensuring that when a source correction is committed, it flows through the same pipeline path the original mutation traveled.

The semantic fingerprint in detail

The semantic fingerprint deserves special attention because it bridges all three architectures. A fingerprint is a compact, deterministic representation of a dataset's distributional properties that can be compared across runs with a single similarity score rather than running the full battery of statistical tests on every comparison.

A well-designed semantic fingerprint for a financial dataset captures: the mean and standard deviation of each numeric column (normalized), the p5/p50/p95 percentile triplet, the null rate, the distinct value count ratio (distinct / total), the top-5 value frequency distribution for categorical columns, and the cross-column correlation matrix flattened to a vector. These are concatenated and stored as a fixed-length vector per dataset per run.

Fingerprint comparison is then a vector distance computation — cosine similarity or Euclidean distance — rather than a full statistical test battery. A similarity below a calibrated threshold triggers deep-dive analysis. Fingerprints stored as a time series create an efficient anomaly detection surface: you can spot when a dataset's fingerprint trajectory changes direction without running any statistical test at all. This is particularly valuable for retroactive mutation detection — if the fingerprint for a prior-period dataset changes between successive pipeline runs, something has altered historical data, and the change is detectable regardless of how small it is.

The practical implication for the solution architecture: fingerprint computation and comparison can run in milliseconds on modern data platforms. It is cheap enough to run on every dataset at every pipeline boundary, providing a universal first-pass anomaly detector that triggers the more expensive statistical analysis only when a potential anomaly is detected.

The Profiling Sidecar: Architecture and Design in Detail

The profiling sidecar is the most practically deployable and operationally mature pattern for value profiling at scale. Unlike inline profiling — which embeds quality checks directly in the pipeline's critical path and adds latency to every run — the sidecar decouples statistical analysis from data movement entirely. The pipeline produces data at its normal speed. The sidecar observes what was produced, runs its analysis asynchronously, and raises an alarm if something looks wrong. The two systems share nothing except the metadata events that flow between them.

This decoupling is not a cosmetic separation. It changes the economics of profiling fundamentally. When profiling runs inline, every additional check adds to pipeline latency, and pipeline owners face constant pressure to remove checks to meet SLA targets. When profiling runs in a sidecar, checks can be as computationally expensive as needed — Isolation Forest scoring, autoencoder reconstruction, full percentile profile comparison — without anyone caring how long they take, because the pipeline has already completed and delivered its data. The tradeoff is that sidecar detection is slightly delayed relative to inline detection: the pipeline has already written its output by the time the sidecar raises an alert. In practice, for batch pipelines running on hourly or daily schedules, this delay is irrelevant. For streaming pipelines with sub-minute latency requirements, a hybrid approach is needed.

The core design principle: observe, don't intercept

The most important design decision in a profiling sidecar is that it must never sit in the data path. It reads profile events from the metadata bus; it does not read the actual data. This means the sidecar scales independently of data volume — profiling a table with 10 billion rows costs no more than profiling one with 10 million rows, because the sidecar only ever sees the statistical summary, not the underlying records.

The pipeline's responsibility is to compute and emit a profile event at the completion of each stage. A profile event is a structured metadata record containing the column statistics for the dataset produced at that stage: counts, distributions, quality metrics. The sidecar consumes these events, enriches them with baseline comparisons, and produces anomaly assessments. The actual data never flows through the sidecar.

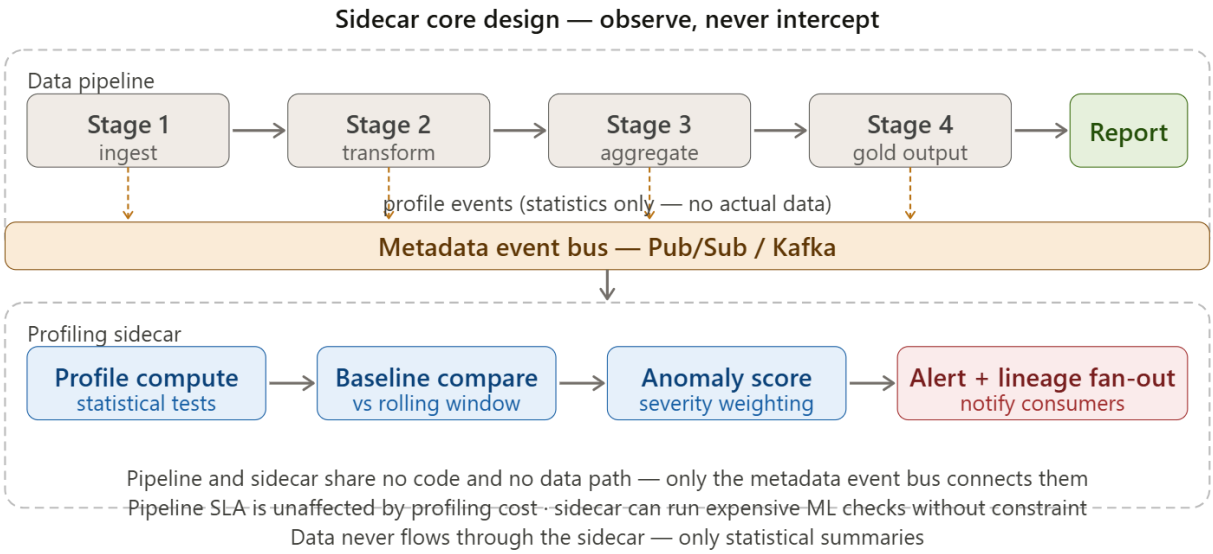

What a profile event contains

Every time a pipeline stage completes, it emits a profile event onto the metadata bus. The profile event is the atomic unit of the sidecar's input stream. Its design determines what the sidecar can and cannot detect — an under-specified profile event produces a sidecar that is blind to important mutation classes.

A well-designed profile event for a financial data pipeline contains four categories of information. The first is the identity payload: which dataset, which pipeline run, which stage, which partition (date, business unit, entity), and the run timestamp. This uniquely identifies the event in the lineage graph and allows the sidecar to retrieve the correct baseline for comparison.

The second category is the structural summary: row count, column count, schema hash (a deterministic hash of the column names and types), and partition boundaries. Schema hash changes are themselves a signal — a silent schema change that preserves column names but alters types is invisible to structural checks but detectable here.

The third category is the statistical payload: for each column, the full statistical profile — mean, standard deviation, min, max, null count, distinct count, and the percentile vector (p1, p5, p10, p25, p50, p75, p90, p95, p99, p99.9). For categorical columns, the top-50 value frequency distribution. For date columns, min date, max date, and gap count.

The fourth category is the cross-column payload: pairwise Spearman correlations for a declared set of monitored column pairs, ratio values for declared ratio invariants, and

join match rates for foreign key relationships. These are not computed for every possible column pair — that would be computationally expensive and mostly useless — but for a declared set of business-critical relationships specified in the data contract.

The five internal engines

The sidecar is not a single process. It is a coordinated set of five engines, each consuming the same profile event stream and producing anomaly signals that feed the correlation layer.

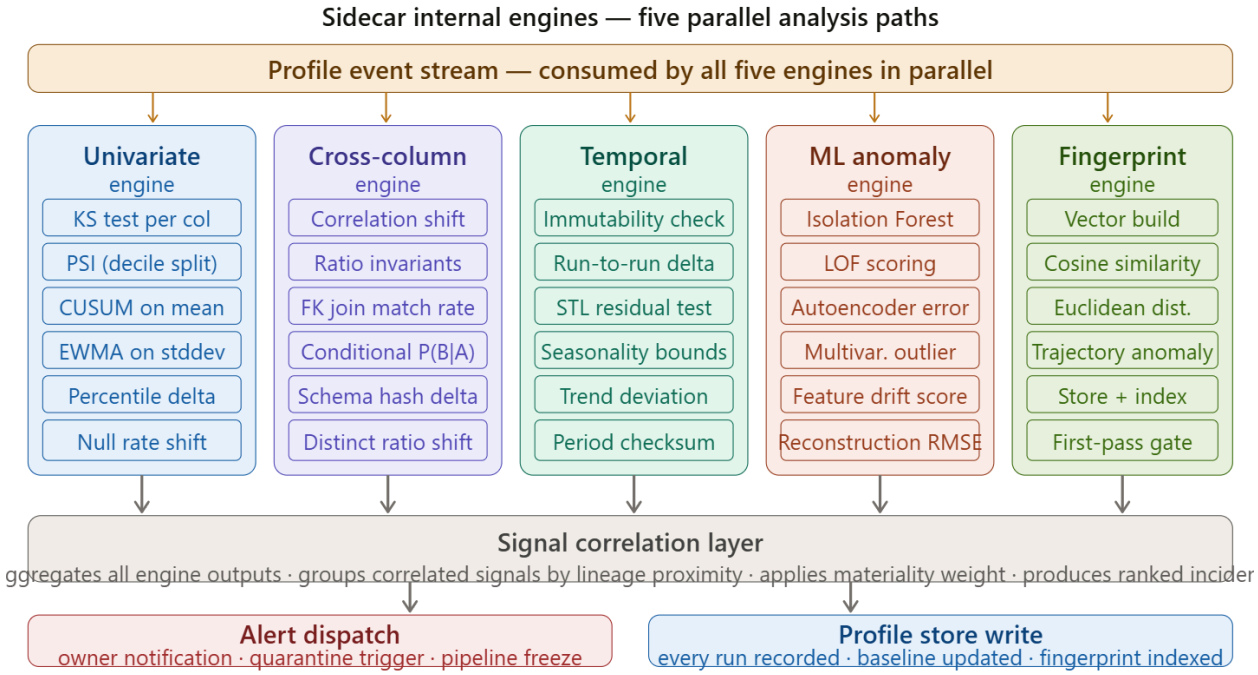

Engine 1 — Univariate engine

This is the fastest and most broadly applicable engine. For every column in every profile event, it runs the KS test against the rolling 30-day baseline distribution, computes the PSI by decile, runs CUSUM to detect sustained mean shifts, and runs EWMA to detect gradual standard deviation drift. It also computes the delta between the current percentile vector and the baseline percentile vector, which catches scale changes and tail

mutations that neither KS nor PSI reliably detect. The null rate shift check is simple but high-value: a null rate that moves from 0.3% to 4.2% in a single run almost always indicates a join failure or a source system problem upstream.

Engine 2 — Cross-column engine

This engine operates on the declared set of monitored column relationships specified in the data contract for each dataset. For a financial fact table, this set might include: the Spearman correlation between interest income and average earning assets, the ratio of total debits to total credits (which should be approximately balanced), the foreign key join match rate between the fact table and the account dimension, and the conditional distribution of transaction amounts given transaction category. The key design point is that this engine only runs checks that have been declared — it does not attempt to profile every possible column pair, which would be both computationally expensive and prone to spurious correlation alerts.

Engine 3 — Temporal engine.

This engine maintains the historical checksum registry and compares it on every run. At the end of each pipeline run, it stores a SHA-256 hash of the key values in each immutable partition. On subsequent runs, it recomputes these hashes and compares. Any change in a prior-period partition hash is a retroactive mutation signal, regardless of how small the change is. The STL residual test separates the seasonal, trend, and residual components of each key metric's time series and applies anomaly detection to the residual, enabling accurate detection of genuine mutations against a background of seasonal variation. The period checksum is particularly important for financial pipelines because it is the only method that reliably detects retroactive restatements.

Engine 4 — ML anomaly engine

This engine runs the computationally expensive methods that would be impractical inline: Isolation Forest scoring on the multi-column profile vector, Local Outlier Factor for density-based anomaly detection, and autoencoder reconstruction error for datasets that have been designated as high-value and warranted the effort of training a reconstruction model. The input to these ML methods is not the raw data but the statistical profile vector — the same vector that the fingerprint engine builds. This means the ML engine is computing anomaly scores on the profile of the data rather than the data itself, which is both computationally efficient and conceptually appropriate: what the engine is measuring is whether the statistical character of the data is anomalous, not

whether individual records are anomalous.

Engine 5 — Fingerprint engine

This engine is the sidecar's universal first-pass detector. It builds the semantic fingerprint vector for each profile event, compares it against the rolling baseline centroid using cosine similarity and Euclidean distance, and checks the fingerprint time series for trajectory anomalies. The fingerprint comparison completes in milliseconds and serves as a gate: if the fingerprint is within normal bounds, the other four engines' results are archived without triggering alerts. If the fingerprint is anomalous, the other engines' results are escalated to the correlation layer for full assessment. This gating behavior is what makes the sidecar computationally tractable at scale — the expensive engines only do their full work on the small fraction of runs where the fingerprint says something may be wrong.

The baseline store: the sidecar's memory

The sidecar's ability to detect mutations depends entirely on the quality of the baseline it compares against. A contaminated baseline — one that includes mutated batches — will normalize the mutation and make future mutations harder to detect. A stale baseline — one that does not reflect legitimate seasonal evolution — will generate constant false positives on natural variation. The baseline store design is therefore one of the most important engineering decisions in the sidecar architecture.

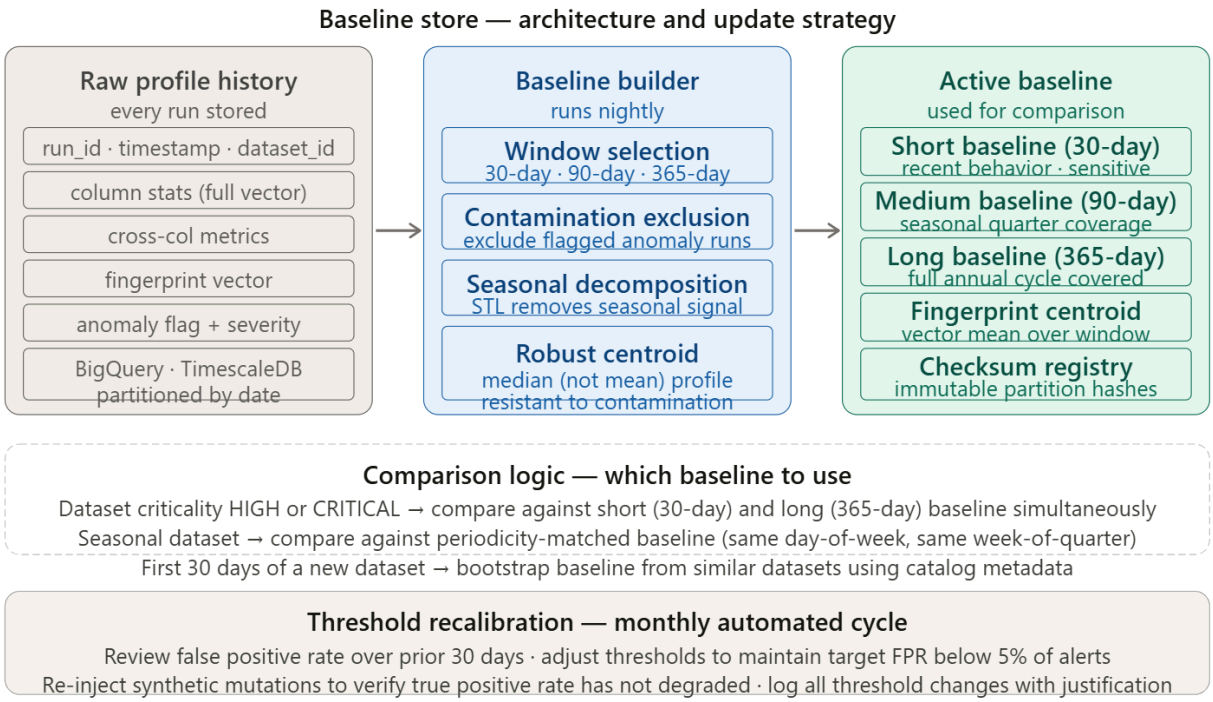

The baseline store maintains three rolling windows simultaneously: 30-day, 90-day, and 365-day. This is not redundancy — each window serves a different detection purpose. The 30-day baseline is the most sensitive, calibrated to recent behavior; it catches mutations quickly but may be triggered by legitimate slow-moving business changes. The 90-day baseline covers a full financial quarter and is the right comparison for datasets with strong quarterly seasonality. The 365-day baseline is the definitive reference for catching annual pattern breaks and long-term definitional drift.

The contamination exclusion step is critical. When the baseline is computed as a simple rolling mean or median over the prior window, any runs that contained genuine mutations are included in that computation and shift the baseline toward the mutated values. Over time, a mutation that was not corrected will be absorbed into the baseline and become undetectable. The exclusion step removes all runs that were flagged as anomalous by the sidecar before computing the baseline centroid. This is only correct if the anomaly flags are accurate — which is why false positive management is so important: an excessively sensitive sidecar that flags normal runs as anomalous will inadvertently thin its own baseline and reduce its future sensitivity.

Using the median rather than the mean as the baseline centroid is a deliberate choice for robustness. The mean is sensitive to outliers in the historical window; the median is not. For financial data that includes rare but legitimate extreme events — a quarter-end

settlement spike, a year-end restatement correction — the mean-based baseline will be pulled toward these events even after they are excluded from the contamination filter, because the exclusion filter operates on flagged runs, not on individual extreme values within clean runs. The median is immune to this effect.

Signal correlation: turning multiple alerts into one incident

A single upstream mutation typically generates many correlated signals across the five engines. Without a correlation layer, this produces an alert storm: the analyst receives one alert for the KS test failure on column A, another for the PSI failure on column A, another for the ratio invariant failure between column A and column B, another for the Isolation Forest anomaly score on the multivariate profile, and another for the fingerprint distance anomaly. All five alerts point to the same mutation, but the analyst must piece that together from five separate notifications.

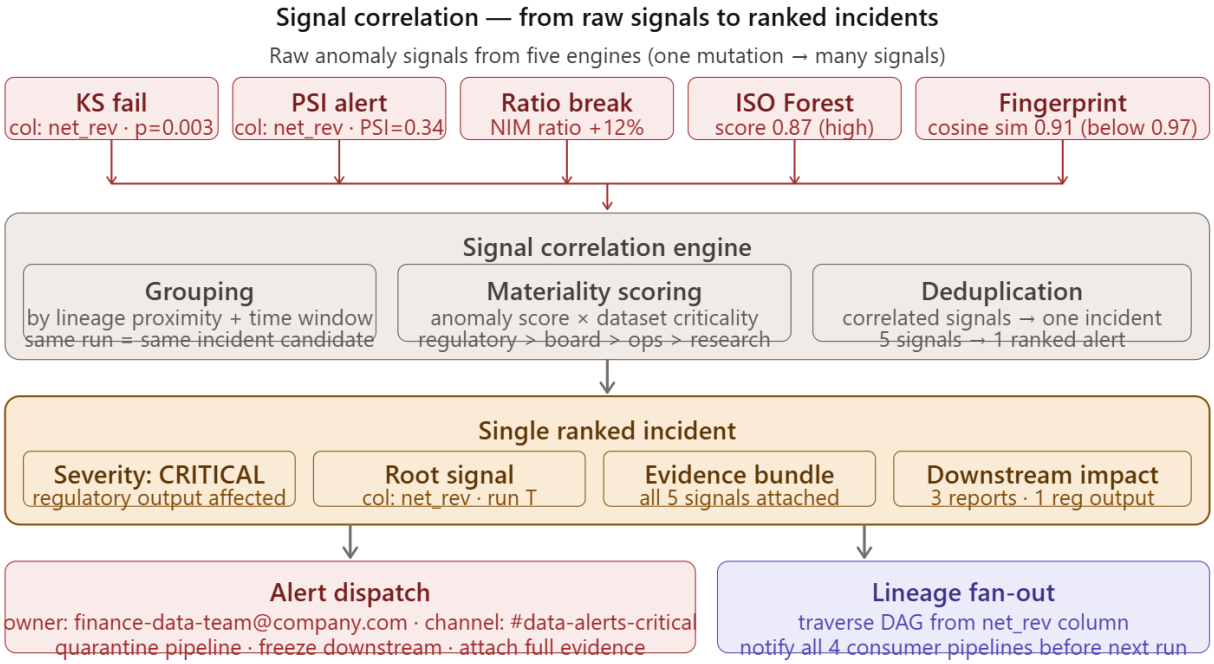

The signal correlation engine groups anomaly signals by two criteria: lineage proximity (signals referencing columns that are in the same lineage neighborhood) and temporal coincidence (signals generated from the same pipeline run or within a configurable time

window). Signals that pass both grouping criteria are merged into a single incident candidate, with all contributing signals attached as evidence.

Materiality scoring then elevates or suppresses the incident based on the downstream criticality of the affected data. An anomaly on a column that feeds only an internal research dashboard scores low regardless of its statistical magnitude. The same anomaly on a column that feeds the CET1 capital ratio scores critical regardless of whether the magnitude seems small. The materiality weights are maintained in the data catalog as attributes of each dataset's downstream consumers.

Deduplication is the final step: five signals about the same column in the same run become one ranked incident with a severity, a root signal, a full evidence bundle, and a downstream impact assessment. The analyst receives one notification, not five.

The lineage fan-out: why proactive notification matters

Most data quality systems are reactive — they alert when something has failed. The lineage fan-out makes the sidecar proactive. When an anomaly is detected on a dataset at pipeline stage N, the sidecar traverses the downstream lineage graph from that point and identifies every consumer pipeline, every derived dataset, and every report that depends on the affected data. It notifies those consumers immediately — before their next scheduled run — with the message that the data they depend on has been flagged, the nature of the anomaly, and the recommended action.

This matters because the time between an upstream anomaly and a downstream consequence is not zero. A mutation detected in a morning pipeline run might not surface as a wrong number in a report until the afternoon. Without proactive fan-out, the analyst who runs the afternoon report discovers the wrong number and begins investigating from scratch. With proactive fan-out, they have already been notified that the upstream data was anomalous three hours ago, the anomaly is already being investigated, and they should hold their report pending resolution.

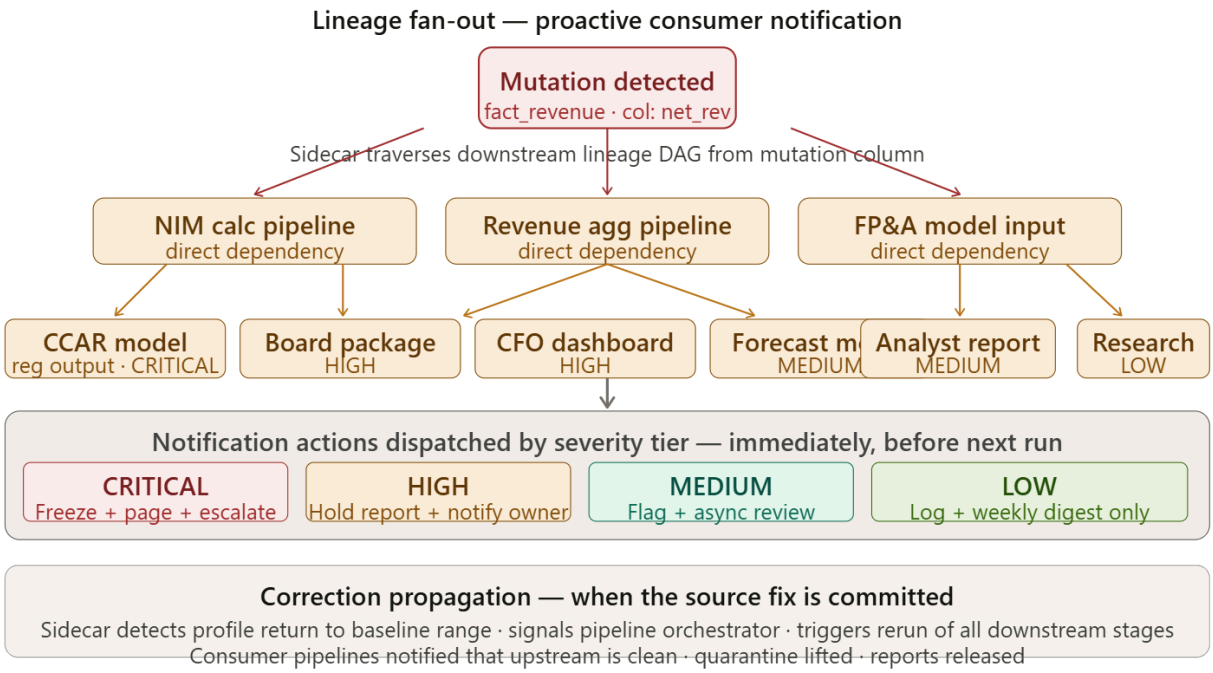

Operational characteristics and deployment considerations

Latency. The sidecar's detection latency is the time from pipeline completion to alert dispatch. For a batch pipeline, this is typically 2–5 minutes: the time for the profile event to arrive on the bus, be consumed by the sidecar's engines, pass through the correlation layer, and trigger the alert. For high-criticality datasets, the sidecar can be configured to run in a pre-write interception mode, where the pipeline holds its final write until the sidecar clears the pre-write profile — essentially a hybrid inline/sidecar pattern for the final stage only.

Throughput. A sidecar implementation using Dataflow (GCP) or Flink (open source) can process thousands of profile events per second without difficulty, because profile events are small (typically 10–50 KB per event) and the statistical computations are embarrassingly parallel. The ML anomaly engine is the only bottleneck, and it is addressed by configuring ML scoring only for datasets above a designated criticality threshold.

Cold start. A new dataset has no historical baseline. The cold start problem is addressed by bootstrapping: the catalog metadata is used to identify similar datasets (same domain, same schema pattern, same business owner), and their baselines are used as a starting prior for the new dataset. After 30 days of clean runs, the dataset's own baseline is

mature enough to replace the bootstrapped prior.

Failure handling. The sidecar must be designed so that its failure does not fail the pipeline. The pipeline runs regardless of the sidecar's state. If the sidecar is unavailable, profile events accumulate on the bus and are processed when it recovers. The worst case of a sidecar outage is a gap in profiling coverage for the duration of the outage — not a pipeline failure. This separation of concerns is one of the architectural advantages of the sidecar pattern over inline profiling.

Profile event schema evolution. As pipelines evolve and new columns are added, the profile event schema must accommodate new columns without breaking existing baseline comparisons. The recommended approach is to version the profile event schema and maintain the baseline per schema version, rebuilding the baseline when a schema version change is detected. The schema hash in the profile event's identity payload serves as the version identifier.

Conclusion:

The transition to a state of **Data Nirvana** requires more than just moving data to the cloud; it requires a fundamental reimagining of how the enterprise verifies the integrity of its most critical asset. The **Value Profiling Sidecar** represents this paradigm shift—moving away from the brittle, high-latency manual checks of the past and toward an autonomous, observability-first future.

By decoupling complex statistical profiling from the execution path, we eliminate the "latency tax" that has historically hindered enterprise agility. This architecture does more than just flag errors; it establishes a **Semantic Fingerprint** for every data product, providing the high-fidelity context necessary for AI agents to operate with total confidence. For the C-Suite, the Sidecar is the engine that converts "dark data" into a "bankable" strategic asset, effectively shifting the data organization from a defensive cost center into a proactive growth engine.

In an era where autonomous workflows are beginning to make real-time financial and operational decisions, "good enough" data quality is no longer a viable strategy. The Sidecar provides the immutable trust layer required for the **Cognitive Enterprise**. The framework is defined, the statistical methods are verified, and the architecture is ready for enterprise-scale deployment. The path to zero-latency intelligence is now clear:

automate the stewardship, secure the statistical baseline, and empower the business to move at the speed of curiosity.